

Components for Sorting Network Proofs

A catalogue of the constructions used in the verification of djbsort

Laurent Théry

Abstract. This note catalogues the constructions used in the two verifications of djbsort — the portable `int32` sort and the AVX2 sort — in the style of *A Formalisation of Algorithms for Sorting Network*. It records what each component is, why it exists, and which of them are shared. Components introduced since that note are marked **NEW**.

1 Introduction

A network is a sequence of connectors acting on m lines. The original note develops that model and three algorithms that build sorting networks: the bitonic sorter, Batcher’s odd-even merge sorter, and Knuth’s odd-even exchange sorter. This note starts where that one stops.

Verifying real code needed more than the model. Two programs were verified:

- `int32` — djbsort’s portable `sort.c`, which implements Knuth’s merge exchange (TAOCP 5.2.2M) as a flat loop nest;
- `avx2` — djbsort’s vectorised sort, whose 8×8 transpose and sign-flip masks realise a periodic bitonic network.

Each raised a different difficulty, and each grew its own machinery. Section 2 recalls the foundation. Section 3 describes the components that are generic and now shared. Sections 4 and 5 describe what remains specific to each track. Section 6 summarises status and file layout.

The organising observation is that the foundation offers only **binary** combinators — everything is built by splitting a network in halves. Both verifications needed to say instead “run this small network **everywhere**”, on blocks or on residue classes, and both had to build that themselves.

2 The foundation

2.1 Connectors and networks

A connector links each line to a partner, involutively, and records for each line whether the comparison is flipped.

```
Record connector (m : nat) := connector_of {
  clink : {ffun 'I_m -> 'I_m};
  cflip : {ffun 'I_m -> bool};
  cfinv : [forall i, clink (clink i) == i];
  cflipinv : [forall i, cflip (clink i) == cflip i] }.

Definition cfun c t :=
  [tuple let min := min (tnth t i) (tnth t (clink c i)) in
    let max := max (tnth t i) (tnth t (clink c i)) in
    if i <= clink c i then if cflip c i then max else min
    else if cflip c i then min else max | i < m].

Definition network := seq (connector m).
Definition nfun n t := foldl (fun t c => cfun c t) t n.
Definition sorting := [qualify n | [forall r : m.-tuple bool, sorted <=%0 (nfun n r)]].
```

A connector is one **parallel stage**: it performs several disjoint comparators at once. This is worth stressing, because it is exactly where the `int32` track ran into trouble (§4.2).

2.2 Building networks

The combinators available are all binary — they split $m + m$ lines into two halves, either contiguous or by parity.

<code>cmerge, cdup</code>	contiguous halves, at connector level
<code>ceomerge, ceodup</code>	even/odd halves, at connector level
<code>nmerge, ndup</code>	contiguous halves, at network level
<code>neomerge, neodup</code>	even/odd halves, at network level
<code>cswap i j</code>	a single comparator between lines i and j
<code>codd_jump k</code>	odd lines linked to their k -th successor (k odd)
<code>ceswap</code>	even lines linked to their odd neighbour

The semantics of the parity split is the one identity everything else rests on:

```
Lemma nfun_eodup (n : network m) (t : (m + m).-tuple A) :
  nfun (neodup n) t = teocat (nfun n (tetake t)) (nfun n (totake t)).
```

2.3 The three sorters

```
Fixpoint half_cleaner_rec b m : network (`2^ m) := ... (* bitonic *)
Fixpoint bsort m : network (`2^ m) := ...
Fixpoint bfsort (b : bool) m : network (`2^ m) := ...

Definition batcher_merge {m} : connector m := codd_jump 1. (* odd-even merge *)
Fixpoint batcher m : network (`2^ m) :=
  if m is m1.+1 then ndup (batcher m1) ++ batcher_merge_rec m1.+1 else [::].

Fixpoint knuth_jump_rec m k r : network m :=
  if k is k1.+1 then codd_jump r :: knuth_jump_rec m k1 (uphalf r).-1 else [::].
Fixpoint knuth_exchange m : network (`2^ m) :=
  if m is m1.+1 then
    neodup (knuth_exchange m1) ++ ceswap :: knuth_jump_rec (`2^ m) m1 ((`2^
m1).-1)
  else [::].
```

with `sorting_batcher` and `sorting_knuth_exchange` proved. Note that “Batcher” names two different networks here; `djbsort`’s portable `sort.c` implements the **exchange** sorter, not the **merge** sorter.

Alongside the recursive `knuth_exchange`, the development also carries an **imperative** version, `iknuth_exchange`, a `seq`-level program with three nested loops `iter1/iter2/iter3` performing swaps. The two are proved to sort independently of one another; nothing connects them. Which of the two a proof picks up turns out to determine its whole shape (§4).

3 Shared algebra NEW

These are generic: they mention no program. They live in `code/common/nalgebra.v` and are used by both tracks.

3.1 Index pairs as networks

A program emits a **list of comparators**. To reason about it as a network:

```

Definition oconn (n : nat) (ab : nat * nat) : option (connector n) :=
  obind (fun i => omap (fun j => cswap i j) (insub ab.2)) (insub ab.1).
Definition pnet (n : nat) (ps : seq (nat * nat)) : network n := pmap (oconn n) ps.

Lemma pnet_cons (n x y : nat) ps (xn : x < n) (yn : y < n) :
  pnet n ((x, y) :: ps) = cswap (Sub x xn) (Sub y yn) :: pnet n ps.

```

`pnet` produces one connector per **single** comparator, so it always yields a fully sequential network. Recovering parallel stages from it is what §3.4 is about.

3.2 Flip-free connectors

```

Definition cnoflip n (c : connector n) : bool := [forall i, ~~ cflip c i].
Definition nnoflip n (nt : network n) : bool := all (@cnoflip n) nt.

```

This is the second conjunct of the note’s `ctransp`, without its requirement that links be $i \pm 1$ — a requirement `codd_jump r` violates for $r > 1$. Closure lemmas: `cnoflip_odd_jump`, `cnoflip_eomerge`, `nnoflip_neodup`, `nnoflip_neotile`.

3.3 Commutation

```

Definition cdisjoint m (c1 c2 : connector m) : Prop :=
  forall i : 'I_m, (clink c1 i == i) || (clink c2 i == i).

Lemma cfun_comm m (c1 c2 : connector m) t :
  cdisjoint c1 c2 -> cfun c1 (cfun c2 t) = cfun c2 (cfun c1 t).
Lemma nfun_nswap m (n1 n2 : network m) (c1 c2 : connector m) t :
  cdisjoint c1 c2 ->
  nfun (n1 ++ c1 :: c2 :: n2) t = nfun (n1 ++ c2 :: c1 :: n2) t.

```

Two connectors that share no line commute, so adjacent disjoint connectors in a network may be exchanged. This is what licenses reordering a comparator sequence.

3.4 Reading a connector back as comparators

The inverse of `pnet`: which comparators does a connector perform?

```

Definition cpairs n (c : connector n) : seq (nat * nat) :=
  pmap (fun i : 'I_n =>
    if (i < clink c i)%N then Some (nat_of_ord i, nat_of_ord (clink c i))
    else None)
  (enum 'I_n).
Definition nstages n (nt : network n) : seq (nat * nat) :=
  flatten (map (@cpairs n) nt).

Lemma nfun_pnet_cpairs n (c : connector n) t :
  cnoflip c -> nfun (pnet n (cpairs c)) t = cfun c t.
Lemma nfun_pnet_nstages n (nt : network n) t :
  nnoflip nt -> nfun (pnet n (nstages nt)) t = nfun nt t.

```

`nfun_pnet_cpairs` is the bridge between the sequential and the parallel view: a connector’s comparators are pairwise disjoint, because `clink` is an involution, so performing them one at a time is the same as performing the stage. It is what makes a flat program trace comparable with a structured network at all.

The proof is worth a remark, because the obvious route is unpleasant. Case splitting on a line and cutting `enum 'I_n` around the comparator that touches it degenerates into list surgery. Instead one generalises over the list l of **generating** lines — those j with $j < \text{clink } cj$, which is what `cpairs` scans the enumeration for — and describes the whole run in closed form:

```
tnth (nfun (pnet n [seq (j, clink c j) | j <- l]) u) i
= if i \in l          then min (u i) (u (clink c i))
  else if clink c i \in l then max (u (clink c i)) (u i)
  else                u i
```

Each induction step then only checks the three positions the head comparator can reach. The fact that makes this work is that a generator's partner is never itself a generator.

Also here: `cpairs_bounded`, and `tnth_nfun_pnet_avoid` (a line touched by none of the comparators keeps its value).

3.5 Iterated deinterleave

`neodup` distributes over concatenation:

```
Lemma neodup_cat n (n1 n2 : network n) :
  neodup (n1 ++ n2) = neodup n1 ++ neodup n2.
```

which lets a recursive network be flattened. Iterating `neodup` naively is painful, though: it goes `network m -> network (m + m)`, so j iterations land in a tower $(m + m) + (m + m) + \dots$ and every equation drowns in casts. Since `e2n` is defined by doubling rather than by `expn`, the type 2^{j+1} is **definitionally** $2^j + 2^j$, so indexing the iteration by the exponent keeps the type a power and no cast is ever needed:

```
Fixpoint neotile q (net : network (`2^ q)) j : network (`2^ (j + q)) :=
  if j is j1.+1 then neodup (neotile net j1) else net.
```

with `neotile0`, `neotile5`, `neotile_cat`, `nnoflip_neotile`. This is the interleaved sibling of the AVX2 track's blocked `ntile` (§5.4); the two solve the same problem for the two ways of splitting an array.

Finally, the index-level counterpart of `nfun_eodup`, saying that the enumeration of `'I_(m + m)` is the sub-problem's enumeration with each line a expanded to the adjacent pair $2a, 2a + 1$:

```
Lemma iota_eocat m : iota 0 (m + m) = flatten [seq [:: a.*2; a.*2.+1] | a <- iota 0 m].
Lemma enum_ord_eocat m :
  enum 'I_(m + m) = flatten [seq [:: elift a; olift a] | a <- enum 'I_m].
```

3.6 Comparators under deinterleaving NEW

Combining the previous two sections: deinterleaving doubles every comparator, running (a, b) on the even lines and on the odd lines, as $(2a, 2b)$ and $(2a + 1, 2b + 1)$. Because the enumeration visits $2a$ immediately before $2a + 1$, the two copies come out **adjacent**, so the comparator list of a deinterleaved network is the original list with each entry expanded in place.

```
Definition pdup (ps : seq (nat * nat)) : seq (nat * nat) :=
  flatten [seq [:: (ab.1.*2, ab.2.*2); (ab.1.*2.+1, ab.2.*2.+1)] | ab <- ps].
Fixpoint pdupn j ps := if j is j1.+1 then pdup (pdupn j1 ps) else ps.
```

```

Lemma cpairs_eodup n (c : connector n) : cpairs (ceodup c) = pdup (cpairs c).
Lemma nstages_neodup n (nt : network n) : nstages (neodup nt) = pdup (nstages nt).
Lemma nstages_neotile q (net : network (2^q)) j :
  nstages (neotile net j) = pdupn j (nstages net).

```

with `pdup_cat`, `pdup_flatten`, `pdup_pmap`, `neodupE` (`neodup nt is map ceodup nt`) and `clink_ceodup_e / _o`.

Two groups of small generic lemmas support this and §4.5. First, pushing `map`, `filter` and `pmap` through `flatten` in an explicit form — `map_flatten_seq`, `filter_flatten_seq`, `pmap_flatten_seq`, `flatten_map_filter`. The library’s `map_flatten` reassociates into the `[seq _ | x <- s, y <- x]` notation, which `map_comp` can then no longer see through, so the explicit versions are what let two flattened comprehensions be compared termwise. Second, two arithmetic facts,

```

Lemma divn_double a p : 0 < p -> a.*2 %/ p.*2 = a %/ p.
Lemma divn_doubleS a p : 0 < p -> a.*2.+1 %/ p.*2 = a %/ p.

```

which say that doubling numerator and denominator leaves a quotient alone, with or without the odd offset the odd lines carry. That is exactly what makes a test of the form “bit k of the index is clear” survive deinterleaving.

4 The `int32` track

`sort.c` implements Knuth 5.2.2M as a flat loop: `p` descends from `top` by halving, and for each base position the whole distance cascade is run at once.

4.1 The comparator model

The `C` is transcribed as the exact list of comparators it emits, in order.

```

Definition me_top (n : nat) : nat := top_loop n 1 n. (* sort.c's `top` *)
Definition level_pairs (N p d : nat) (b : bool) : seq (nat * nat) :=
  [seq (i, i + d) | i <- [seq i <- iota 0 N | (i + d < N) && (odd (i %/ p) ==
b)]]].
Definition casc_pairs (N top p : nat) : seq (nat * nat) :=
  flatten [seq [seq (j + p, j + r) | r <- [seq r <- halves top top | (p < r) && (j
+ r < N)]]
          | j <- [seq j <- iota 0 N | ~ odd (j %/ p)]]].
Definition me_pairs (n : nat) : seq (nat * nat) :=
  flatten [seq level_pairs n p p false ++ casc_pairs n (me_top n) p | p <- halves
(me_top n) (me_top n)].

Definition int32_sort_network (n : nat) : network n := pnet n (me_pairs n).

```

The cascade is grouped **by position, then by distance** — not by distance, as a transcription of step M3 would naturally give. Reproducing that exact order is what makes `me_pairs n` equal to the trace of the `C`, and it is also the source of every difficulty below.

4.2 Arbitrary n

Three facts reduce the general case to a power of two: `me_pairs_bounded` (every emitted pair is in range), `me_pairs_prune` (`me_pairs n` is `me_pairs` at the next power of two, filtered to the lines below n), and

```

Lemma sorting_pnet_prune (N n : nat) (ps : seq (nat * nat)) :
  n <= N -> all (fun ab => ab.1 < ab.2 < N) ps ->
  pnet N ps \is sorting -> pnet n [seq ab <- ps | ab.2 < n] \is sorting.

```

the last being a general 0-1-principle statement: pruning a sorting pair-network to its low lines still sorts.

4.3 The reification layer

For the power-of-two case the proof leaves the network world entirely. It turns the network into a `foldl` of `seq`-level swaps and matches it against the **imperative** `iknuth_exchange`:

```

Lemma foldl_swap_me_pairs_iknuth s : foldl swap s (me_pairs (size s)) =
  iknuth_exchange s.

```

Everything then happens on `seq (nat * nat)`: `cdep` (two comparators share a line), `indep_blocks`, `is_size_ordered`, `iso`, `dcasc_aux` (the distance-major cascade), and the crux

```

Lemma swseq_casc_dcasc : ... (* distance-major -> position-major transpose *)

```

`cdep` is the `seq`-level shadow of `cdisjoint` (§3.3). The reason the proof took this shape is recorded in the development: `iknuth_exchange` “is the SAME iterative algorithm as `sort.c`, so it matches `me_pairs` directly (unlike the recursive `knuth_exchange`)”. The price is that no part of the argument is reusable, and the AVX2 track could take nothing from it.

4.4 Results

```

Theorem sorting_int32_sort_network n : int32_sort_network n \is sorting.

```

closed under the global context. What is **not** proved is that the C really emits `me_pairs n`; that is stated as an explicit assumption:

```

Parameter sortc_trace : nat -> seq (nat * nat).
Axiom sortc_faithful : forall n, sortc_trace n = me_pairs n.

```

4.5 The algebraic route NEW

`int32_algebraic.v` redoes the power-of-two case in the style of the AVX2 track: stay inside `network`, and prove an equation between networks.

```

Theorem nfun_int32_knuth m t :
  nfun (int32_sort_network (2^m)) t = nfun (knuth_exchange m) t.
Corollary sorting_int32_sort_network_e2n_alg m : int32_sort_network (2^m) \is
  sorting.

```

Sorting then comes from `sorting_knuth_exchange`, and `iknuth_exchange`, `iter1/iter2/iter3` and `swseq_casc_dcasc` all leave the trust path. The remaining step OPEN is

```

Lemma nfun_me_pairs_knuth m t :
  nfun (pnet (2^m) (me_pairs (2^m))) t
    = nfun (pnet (2^m) (nstages (knuth_exchange m))) t.

```

The block **order** is not the obstacle it was believed to be. Pushing `neodup_cat` through the unfolded recursion gives

$$\text{knuth_exchange}(m) = \text{neodup}^{m-1} \text{merge}_1 ++ \dots ++ \text{merge}_m$$

where merge_k is `ceswap` followed by the `knuth_jump_rec` chain on 2^k lines. Each `neodup` doubles distances, so the blocks come out in **decreasing** distance — exactly the order the flat sweep visits p in.

That is the recursive side; §3.6 turns it into a statement about comparator lists, since `nstages` (`neodup nt`) is `pdup (nstages nt)`. The flat side needs the matching claim: that `me_pairs` at 2^{m+1} is `pdup` of `me_pairs` at 2^m , followed by the $p = 1$ block. Its first half is proved,

```
Lemma level_pairs_double N p : 0 < p ->
  level_pairs N.*2 p.*2 p.*2 false = pdup (level_pairs N p p false).
```

— line i of the big problem is $2a$ or $2a + 1$ for a line a of the small one, and both pass the base-pass test exactly when a does, by `ltn_double` for the distance and by `divn_double` / `divn_doubleS` for the bit test.

What remains **OPEN** is the same doubling law for `casc_pairs`, and then the $p = 1$ block against the merge stage. The latter is the old crux, unchanged: `sort.c` emits the cascade position-major and the network emits it distance-major. It is now a statement about networks, so `cfun_comm` applies to it, rather than one about sequences.

5 The `avx2` track

Here the program is not a list of comparators at all: it is a loop nest over a vector width, in which a comparator’s **distance** decides whether it becomes a vector min/max or a lane shuffle. The network is the same for every width, so the proof goes straight to the network.

5.1 Transposition

```
Definition trp (i : 'I_(m * m)) : 'I_(m * m) := ... (* index transpose *)
Definition ttr (t : (m * m).-tuple A) : (m * m).-tuple A := ...
Definition rsh (t : (m * m).-tuple A) : m.-tuple (m.-tuple A) := ... (* reshape *)
Definition fla (M : m.-tuple (m.-tuple A)) : (m * m).-tuple A := ...
Definition cconj (c : connector (m * m)) : connector (m * m) := ...
Definition nttr (net : network (m * m)) : network (m * m) := map cconj net.

Lemma nfun_nttr net t : nfun (nttr net) t = ttr (nfun net (ttr t)).
```

Conjugating a network by the transpose: this is what turns “sort the columns” into “transpose, sort the rows, transpose back”, which is what the AVX2 code actually does.

5.2 Rows and columns

```
Definition crow (c : connector m) : connector (m * m) := ...
Definition nrows (net : network m) : network (m * m) := map crow net.
Definition ncols (net : network m) : network (m * m) := nttr (nrows net).
```

with `nfun_nrows`, `nfun_ncols`, `nrows_sorted`, `ncols_sorted`.

5.3 Sign flips

The code obtains descending order by exclusive-oring with a mask of 0/−1 words rather than by using a different comparator.

```
Definition tflip (msk : (m * m).-tuple bool) (t : (m * m).-tuple A) := ...
Definition ctflip (msk : ...) (c : connector (m * m)) : connector (m * m) := ...
Definition ntflip (msk : ...) (N : network (m * m)) : network (m * m) := ...
Definition mask_luni (msk : (m * m).-tuple bool) : Prop := ... (* lane-uniform *)

Lemma nfun_conj msk cc_net cw_net : ...
Lemma nfun_ntflip_conj msk cc_net : ...
```

Conjugation by a sign flip is the second of the two conjugations the track needs; together with `nttr` they are what let one uniform network stand for the code’s asymmetric-looking phases.

5.4 Tiling and reshape

This is the group that has no counterpart in the foundation, and the one the `int32` track turned out to want too.

```
Fixpoint ntile (net : network (`2^ q)) j : network (`2^ (j + q)) := ...
Definition arsh j (t : (`2^ (j + q)).-tuple A) : (`2^ j).-tuple ((`2^ q).-tuple A) := ...
Definition aflat j (M : (`2^ j).-tuple ((`2^ q).-tuple A)) : (`2^ (j + q)).-tuple A := ...
Definition sqpow : network (`2^ (q + q)) := ...
Definition sqcast (u : (`2^ (q + q)).-tuple A) : (`2^ q * (`2^ q)).-tuple A := ...

Lemma nfun_ntile_arsh net j t : ... (* blockwise semantics *)
Lemma ntile_ntile p net q j : ... (* nested tiling collapses *)
Lemma nfun_tile_sqpow_flat j t : ...
```

`ntile` runs a block network on every block of a larger array, and, exactly as for `neotile` (§3.5), is indexed by the exponent so that no cast appears. `ntile_ntile` collapses nested tilings — the operation that converts a doubly nested loop into a single flat one.

5.5 The sort itself, and its target

```
Definition tmerge_avx2 (b : bool) m (t : (`2^ m).-tuple A) := ... (* the code *)
Fixpoint tsort (b : bool) k : (`2^ k).-tuple A -> (`2^ k).-tuple A := ...
Fixpoint pbsort (b : bool) k : network (`2^ k) := ... (* the target *)

Lemma sorting_pbsort k : pbsort false k \is sorting.
Lemma tsortE b k t : tsort b k t = nfun (pbsort b k) t.
Lemma tmerge_avx2P b m t : ...
```

The target is the **periodic** bitonic network `pbsort`, not the plain `bfsort`; identifying it correctly was itself a step in the development. Padding to a power of two is handled generically by `nfun_pad`, `nfun_pad_sorted`, `nfun_pad_perm` in `sort_generic.v`.

5.6 Results

```
Theorem tsort_avx2_pbsort k t : tsort tmerge_avx2 false t = nfun (pbsort false k)
t.
Corollary avx2_sort_sorted k t : ...
```

with `avx2_sort_perm`, `avx2_sort_pad_sorted`, `avx2_sort_pad_perm`, all closed under the global context. As for `int32`, that the OCaml/C source really runs this network is not formalised; here it is checked empirically instead of assumed, by `ml/trace_check.ml`, which compares the comparator trace of the source against a transcription of the Rocq definitions for every power-of-two width.

6 Status and layout

6.1 What is proved

Result	Status	Meaning
<code>sorting_int32_sort_network</code>	closed	<code>sort.c</code> 's exact network sorts, every n
<code>sortc_faithful</code>	axiom	the C emits that exact trace
<code>tsort_avx2_pbsort</code>	closed	the AVX2 sort computes <code>pbsort</code>
<code>avx2_sort_sorted, _perm</code>	closed	hence it sorts and permutes
<code>nfun_pnet_cpairs</code>	closed	sequential comparators = parallel stage
<code>nstages_neodup</code>	closed	deinterleaving doubles every comparator
<code>level_pairs_double</code>	closed	the base pass survives deinterleaving
<code>nfun_int32_knuth</code>	modulo OPEN	the algebraic route's capstone
<code>nfun_me_pairs_knuth</code>	OPEN	flat sweep = deinterleaved recursion

The single open statement is `nfun_me_pairs_knuth`; everything else listed is closed under the global context. What it still needs is the doubling law for `casc_pairs` and the cascade reordering inside the $p = 1$ block (§4.5).

6.2 Files

<code>code/common/more_tuple.v</code>	tuple and index helpers, the <code>e2n</code> power notation
<code>code/common/nsort.v</code>	connectors, networks, <code>sorting</code> , binary combinators
<code>code/common/nbitonic.v</code>	half cleaners, <code>bsort</code> , <code>bfsort</code>
<code>code/common/nalgebra.v</code> NEW	the shared algebra of §3
<code>code/portable4/proof/nbjsort.v</code>	<code>knuth_exchange</code> and <code>iknuth_exchange</code>
<code>code/portable4/proof/int32_network.v</code>	<code>me_pairs</code> , the reduction facts
<code>code/portable4/proof/int32_reify.v</code>	the <code>seq</code> -level bridge
<code>code/portable4/proof/int32_sort.v</code>	the final theorem, <code>sortc_faithful</code>
<code>code/portable4/proof/int32_algebraic.v</code> NEW	the algebraic route of §4.5
<code>code/avx2/proof/sort_generic.v</code>	<code>gnet</code> , <code>pbsort</code> , padding
<code>code/avx2/proof/sort_transpose.v</code>	§5 in its entirety

6.3 A remark on what was missing

Both tracks needed the same thing and neither could get it from the foundation: a way to say “run this network everywhere”. The AVX2 track built it for blocks (`ntile`) and got a clean development; the `int32` track did not build it, worked around it on lists, and got a development that reuses nothing. The shared algebra of §3 is that missing layer, and `neotile` is the same idea for the other way of splitting an array.